

Credit Evidence Engine

Evidence that an AI assistant in a regulated credit decision told the human what mattered — computed, not judged.

Sriram Krishnan · Clyde · Luca | Mentor: Tosin Adesuyi | NVIDIA Open Models Codefest · 7 October 2026

The problem: a briefing that reads well and is wrong

One referred loan. The assistant's summary, the judge model's grade, and the real number.

APP000044 · NEMOTRON 3.5 LIGHTNING · OUR GPU · 16 SEP 2026

“...the total monthly debt service works out to approximately £630 per month. As a percentage of the applicant's gross monthly income (£1,327), this results in a TMDS ratio of 41.58%. This exceeds the policy threshold of 40%.”

JUDGE · NEMOTRON 3 ULTRA

2 / 2

Intelligible. Actionable.

CASE FILE · ARITHMETIC

47.5%

£630 ÷ £1,326.83. The stated ratio appears nowhere in the file.

- Banks are putting assistants like this in front of the underwriting referral queue. The underwriter decides from the summary.
- Ten runs of this case, same settings: **ten different ratios, eight of them wrong.** Six cited the applicant's age band as a factor. It has no weight in the decision.
- The judge model gave every one of them full marks. **A model cannot tell you what another model got wrong.**
- Under the EU AI Act this is a high-risk system (Annex III 5(b)). The deployer must show the human could see what mattered. Today, nobody can.

Why existing approaches miss it — and what we do instead

Verifier's Law (Shafiq, Wells Fargo): automate what is hard to do but easy to verify.




Make the right answer **known before the model runs**, and every check becomes a comparison, not an opinion.


	GENERIC LLM EVALUATION (DEEPEVAL, RAGAS, LLM-AS-JUDGE)	CREDIT EVIDENCE ENGINE
Where the right answer comes from	Someone labels it, or a judge model decides	Computed from a scorecard on generated cases, before any model call
What a score means	“Faithfulness 0.83” — a model's opinion, averaged	“The briefing omitted the 40% limit in 3 of 20 cases” — checkable by reading the file
Role of the judge model	Primary grader	Readability only. Never completeness or correctness
Protected characteristics	Not a concept	Decoy fields with zero causal weight, declared in the data; citing one is a failure
Output	A metrics dashboard	A tamper-evident evidence pack, by EU AI Act article, for a supervisor

Ground truth by construction

We generate the loan applications, so we know what drove each decision before the assistant reads the file.



How it works



Tooling: Synthetic Data Designer · NIM · NeMo (LoRA) · NeMo Evaluator (runs and gates the pack) · Python · Apache 2.0

Live: one command, on our GPU

MOCKUP — LIVE ON 7 OCT

```
$ evidence run packs/underwriter-sample --sut lightning-nim --repeats 3 --out runs/2026-10-07
pack      underwriter-sample v1.2  20 items · 3 documents each · sha 9f3a...
sut       nvidia/nemotron-3.5-lightning  http://rtx-3se-05-36:8000/v1  temp 0 · seed 7
judge     nano-judge-lora-v3          http://rtx-3se-05-36:8001/v1  on-prem · teacher: ultra · agreement 0.88

running 20 × 3 ... 60/60  1m 52s

check      pass      fail      audit      note
material_omission  17/20    3/20      2          APP000103 APP000212 APP000488 — 40% limit not stated
numeric_fidelity   9/20    11/20     0          11 briefings state a ratio that is not in the file
decoy_citation     12/20    8/20     0          age_band ×6 · dependants ×3
flip_accuracy      18/20    2/20     1
readability (0-2)   mean 1.85
repeat agreement    54/60 (0.90)
                                     nano judge · 3 cases below 1.5
                                     6 flips, all numeric_fidelity

wrote runs/2026-10-07/  manifest · results.jsonl · 60 transcripts · evidence/
```

Every number on this screen traces to a transcript, and every transcript records model id, endpoint, prompt hash, seed, retrieved sections, tokens and latency.

The evidence pack, and the tamper test MOCKUP

```
# Evidence pack – underwriter-sample v1.2 – run 2026-10-07
```

Article 15 – accuracy, robustness

EVIDENCES

```
numeric_fidelity  9/20. 11 briefings state a ratio absent
                   from the case file. Each listed with the
                   number and the file it should come from.
repeat agreement   0.90 over 3 runs, seed pinned. Serving-stack
                   non-determinism recorded, not claimed away.
```

Article 14 – human oversight

EVIDENCES

```
material_omission 17/20. 3 omit the 40% limit the referral
                   turned on.
decoy_citation     12/20. 6 cite age band – zero weight,
                   protected characteristic.
readability        1.85 · nano judge · agreement 0.88 (Ultra),
                   0.85 (human, 20 blind cases)
```

Article 13 · Article 9

CONTRIBUTES

Articles 10 · 12 · 17 · 43

NOT COVERED

Produced from spec credit_underwriting.yaml (sha ...), scorecard v1.
Integrity: checksums.sha256 · evidence verify runs/2026-10-07

```
$ evidence verify runs/2026-10-07
60 transcripts · manifest · results · 3 evidence files
all checksums match.    OK
```

```
$ sed -i 's/41.58%/40.0%/' runs/.../transcripts/APP000044-r1.j
```

```
$ evidence verify runs/2026-10-07
transcripts/APP000044-r1.json: checksum mismatch
expected 3c1f... got 88a0...    FAIL
```

- The reader is a validator or a supervisor, not a data scientist.
- Claims are explicit: **evidences**, **contributes**, or **not covered**. We name what we do.
- A pack that has been edited fails verification, naming the file.

What we found so far

REAL · 16 SEP · TO BE UPDATED WITH THE FULL RUN

8 / 10

runs of one case stated a debt-service ratio not in the file. The judge passed all ten. Only a numeric check catches it.

6 / 10

cited the applicant's age band as a factor against them. It has zero weight and is a protected characteristic.

10 / 10

distinct outputs at temperature 0 with a fixed seed — on the cloud endpoint and on our own vLLM. Reproducibility must be defined as “same verdict across N runs”.

2–3 s

per briefing on our RTX PRO 6000 node, against 8–150 s on the shared cloud endpoint. Same code; two lines of configuration.

- Retrieval failed on the first run: the question reads like a policy question, so the retriever ranked the policy sections above the data, and the assistant wrote a confident briefing with numbers from nowhere. **The transcript records what was retrieved, so the failure can be attributed.**
- The omission check alone is not enough: a briefing can state the breach and still fabricate the ratio. **No single check is sufficient; the pack runs all of them.**

On our own hardware: the assistant, and a judge we can trust

SAME PACK, TWO PLACES

	NVIDIA BUILD · 15 SEP	OUR NODE · 7 OCT
Assistant	Lightning, cloud	Lightning NIM, 1 GPU
Latency per briefing	8–150 s	2–3 s
Judge	Ultra, cloud	Nano + LoRA, 1 GPU
Data leaving the site	Every case file	None
Configuration changed	two lines in .env — recorded in every transcript	

HOW THE ON-PREM JUDGE WAS MADE, AND CHECKED

TARGETS — RESULTS 3 OCT

	VS ULTRA, EXACT	WITHIN 1	VS LUCA, 20 BLIND
Nano 9B, no fine-tune	0.61	0.88	0.60
Nano 9B + LoRA	0.88	0.97	0.85
Ultra (teacher)	—	—	0.90

- **One judge at run time: Nano, on our GPU.** Ultra is the teacher: it graded 400 briefings once, in the cloud, to make the training set. It is not called during a run.
- **Regulation goes in by retrieval, not weights:** Annex III's date moved to Dec 2027 this year. Swap a document, not a model.
- **The judge never grades correctness.** That is what the checks are for.

Run and gated by NeMo Evaluator: 3 repeats, confidence intervals, a policy file that fails the build if omission drops below threshold. Continuous assessment is the SR 26-2 standard.

Day 1 to Day 5

DAY 1

9 SEP

A proposal

- ✓ Pivot declared: evaluate the assistant, not the model
- ✓ Eight contracts specified
- ✓ Two domains chosen; one deferred

DAY 2

16 SEP

One failure, caught

- ✓ SDD recipe, 700 cases, no leaks
- ✓ 20-case pack with marking keys
- ✓ Omission check; 3 Nemotron models live
- ✓ Lightning NIM on our GPU; the numeric finding

DAY 3

23 SEP

Full battery

- numeric_fidelity, decoy_citation
- Runner: 20 × 3 → transcripts
- Judge with regulation passage
- Ultra labels for the student

30 SEP

Evidence

- Evidence pack by article; verify
- Nano LoRA trained and served
- NeMo Evaluator run with gates
- Agreement study

7 OCT

Proven and released

- Cloud vs on-prem manifest diff
- Fresh clone, one command
- Apache 2.0, public
- This demo

Every commit is dated in the repository; BUILD.md carries a change log per day.

Who uses it, and why it lasts

PUBLIC SECTOR

- **The reader of an evidence pack is a supervisor** — ECB, BaFin, FCA, the Federal Reserve. The pack is written for them.
- The same assistant pattern runs inside **public lenders and schemes**: state development banks, government-guaranteed SME and student loans, housing authorities.
- Annex III 5(b) covers **any** creditworthiness assessment of a person — not only by banks.

VIABILITY

- **Apache 2.0.** One complete sample pack ships with the code.
- **Runs where the data is.** Cloud or on-prem, two lines of configuration; nothing leaves the site.
- **A new domain is a new spec and scorecard.** The checks, the pack and the verifier do not change.
- **Regulations change without retraining.** They are documents in a retrieval index.
- **Runs in a harness a bank may already have** — NeMo Evaluator, with pass/fail gates for continuous assessment.

What we do **not**
claim:

the loan decision itself (verified by the loan book, years later)

compliance (we evidence; a reviewer judges)

fairness across a population

Evidence, not opinion.

The right answer is computed before the model runs. Every check is a comparison. The judge grades readability and nothing else. The pack can be re-verified by anyone, years later.

github.com/sriramarun/open-credit-evidence · Apache 2.0 · **Sriram Krishnan** · **Clyde** · **Luca**

Built with Nemotron 3.5 Lightning, Nemotron 3 Ultra, Nemotron Nano 9B, Nemotron Embed · NIM · NeMo · NeMo Evaluator · Synthetic Data Designer